



# Git

## a peek under the hood

Clara Bennett  
@csojinb

follow along:  
[github.com/csojinb/git-under-the-hood/](https://github.com/csojinb/git-under-the-hood/)

# Git: powerful, but leaky

- Like all abstractions, git leaks
- Difficult to master without a solid mental model
- Fear of losing work is a barrier to learning/experimentation
- Taking advantage of git's data-hoarding tendencies requires understanding of how the data is stored

**Solution:** *Gain leverage by learning some internal mechanics*

A wide-angle photograph of a desert landscape. In the foreground, there are large, smooth sand dunes. In the middle ground, a small, isolated cluster of buildings or trees is visible, surrounded by more desert terrain. The background shows distant mountain ranges under a sky filled with scattered clouds.

**What does git store  
when you commit?**

# Core concept: History as snapshots

- To understand how git stores your commits, it's useful to understand the central "philosophy"
- Git "thinks" about version history as a series of **snapshots**, rather than a series of deltas
- A **snapshot** is a complete copy<sup>1</sup> of the project at a particular point in history

---

<sup>1</sup> Unchanged files are not stored multiple times. And, eventually, git will compress versions of the same file together to save space when necessary, e.g. if you want to push to a remote. But the snapshot still decompresses to a complete project copy.

# Representing changes

- Git does not directly save any **actions** that you took, only the **state**
- Differences are **derived** by comparing snapshots
- Actions are **inferred**
- Example (right): git recognizes the rename because the **file content** is

```
3. Clara@Macaque: ~/code/git-under-the-hood (zsh)
[~/code/git-under-the-hood]$ mv images/{,young-}clara-sunglasses.jpg
[~/code/git-under-the-hood]$ git st * [master]
On branch master
Changes not staged for commit:
  (use "git add/rm <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)
)
          deleted:   images/clara-sunglasses.jpg

Untracked files:
  (use "git add <file>..." to include in what will be committed)

          images/young-clara-sunglasses.jpg

no changes added to commit (use "git add" and/or "git commit -a")
[~/code/git-under-the-hood]$ git add --all; git st * [master]
On branch master
Changes to be committed:
  (use "git reset HEAD <file>..." to unstage)

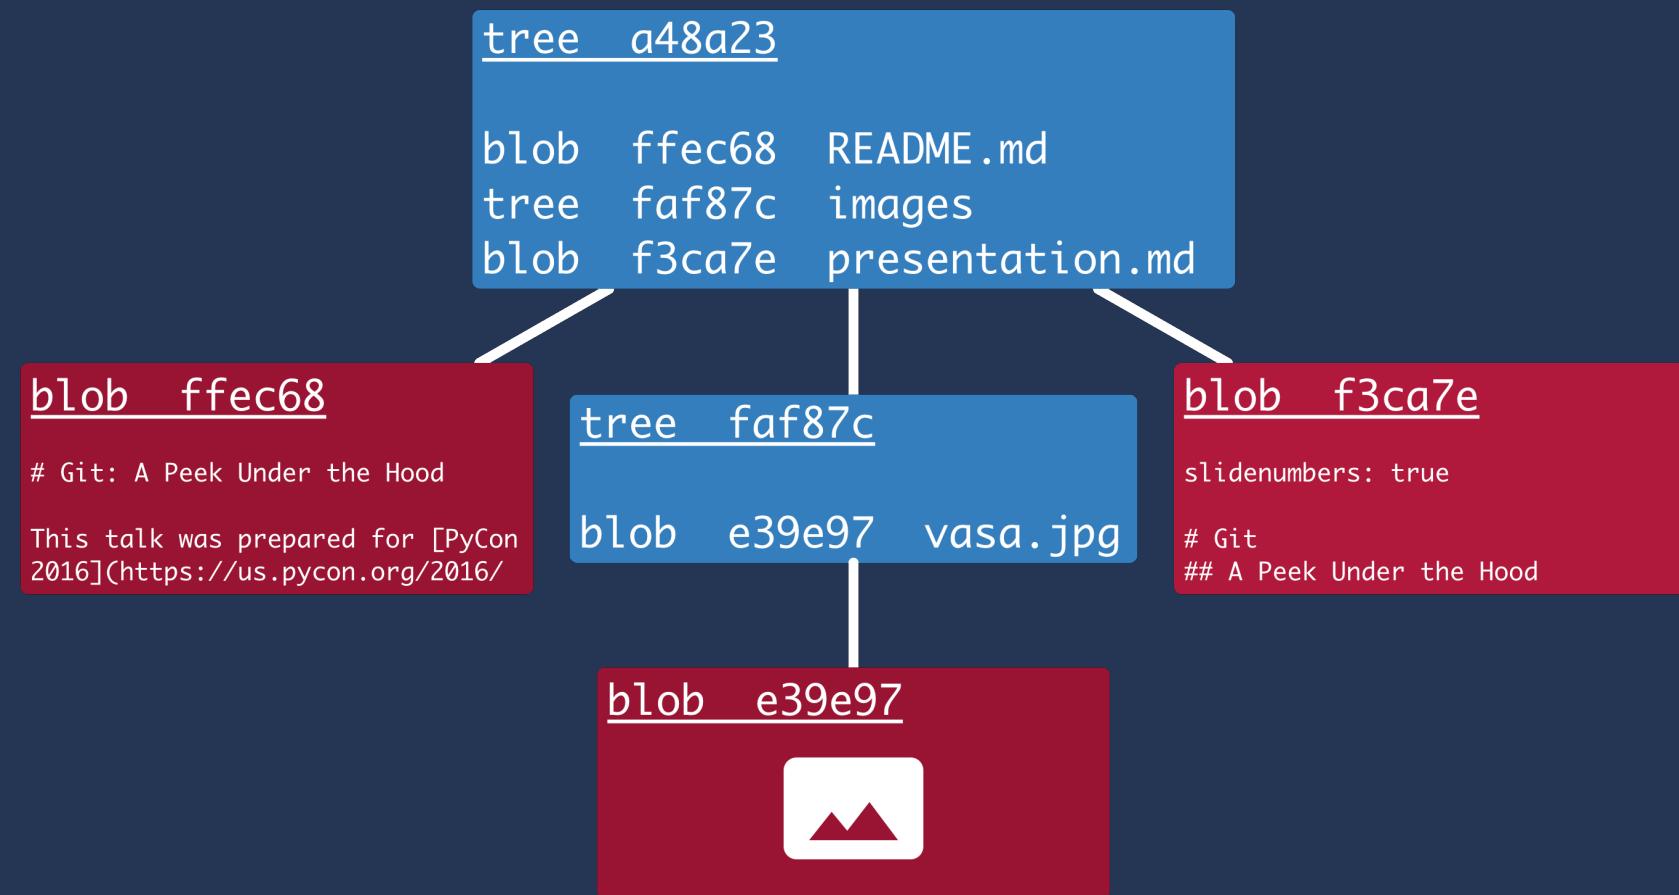
          renamed:   images/clara-sunglasses.jpg -> images/young-clara-sunglasses.jpg

[~/code/git-under-the-hood]$ █
```

<sup>2</sup> Important even if you don't directly use this feature because it affects git's ability to merge intelligently.

# Snapshot storage

- A file snapshot is stored as a text blob, and a directory snapshot is represented as a "tree" object
- Each snapshot is check-summed and stored by SHA-1 value
- Directory trees point to the SHAs of files and directories they contain
- The project snapshot is just the "tree" for the project root directory



# Building a commit

- To make a commit, first you need to stage some changes
- The staging area<sup>3</sup> is just another project snapshot tree
- As changes are staged, new snapshots are created of the affected files/directories, and the staging area is updated
- On commit, the staging area becomes the commit snapshot

---

<sup>3</sup> Sometimes referred to as the "index".

# commit = content + meta-data

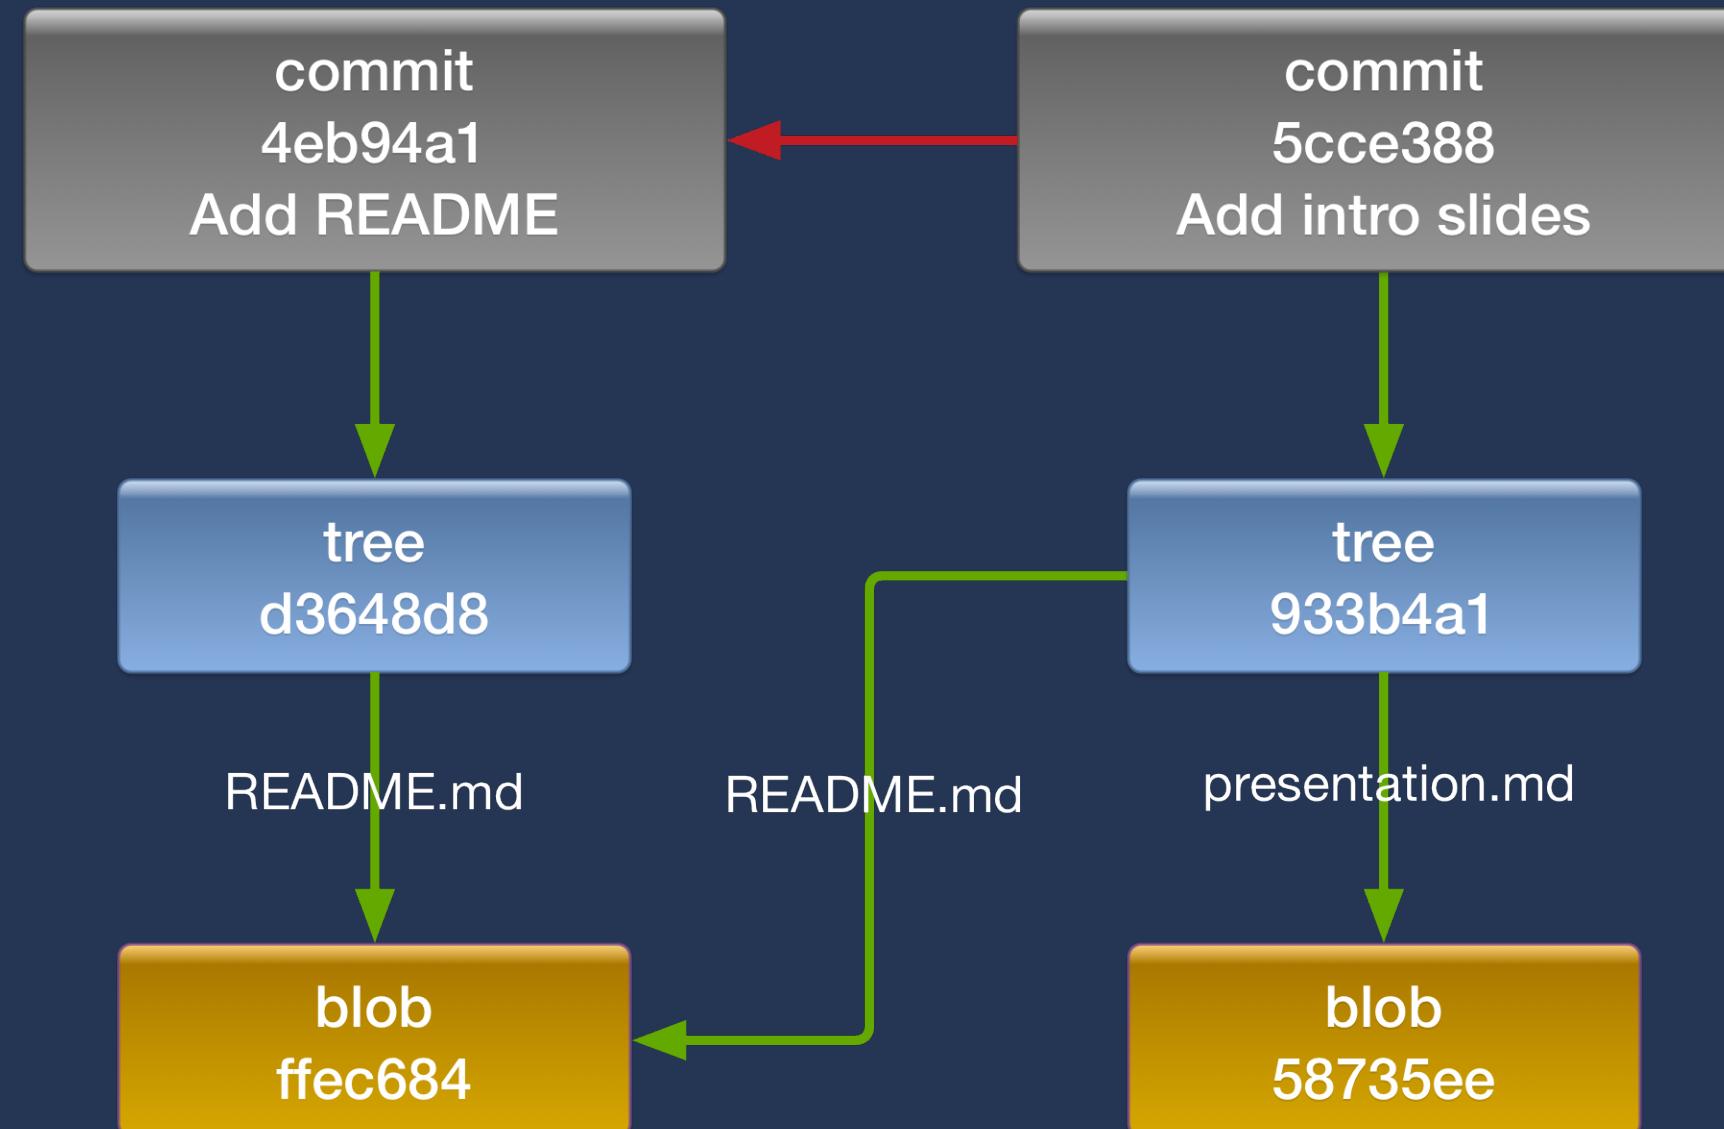
- The final commit object contains a pointer<sup>4</sup> to the **project snapshot** (the content) and some **meta-data**
- The meta-data includes the author, the commit message, and pointer(s) to the **parent commit(s)**<sup>5</sup>
- Note that if either the content or the meta-data is amended, the new commit will have a different SHA checksum value

---

<sup>4</sup> The "pointer" is SHA of the project snapshot

<sup>5</sup> The initial commit has no parents, and merge commits have two or more.

# Visualizing commit storage

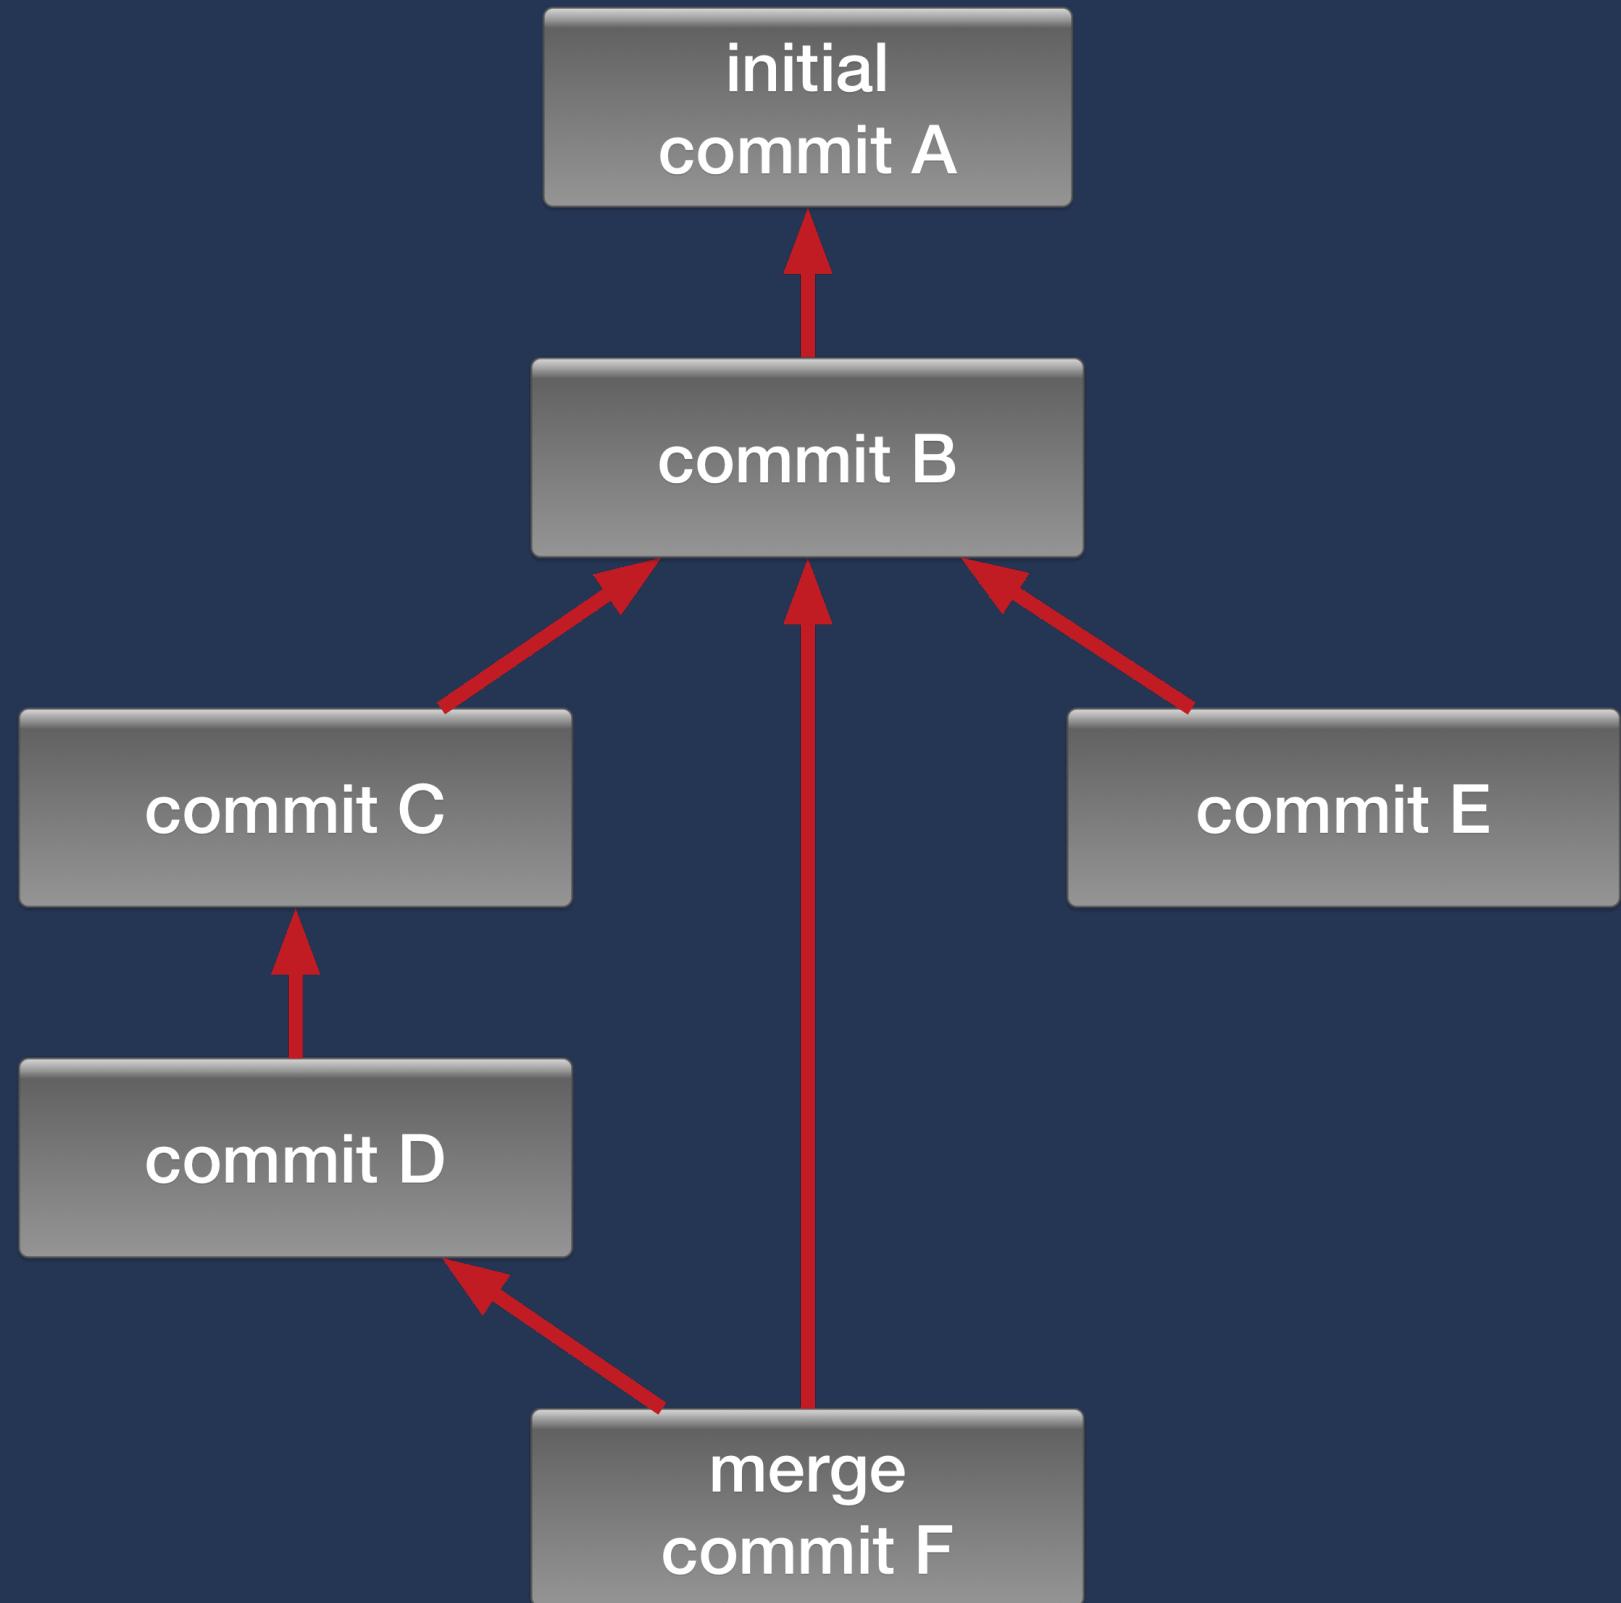


A photograph of a paved path or driveway lined with large, mature trees. The trees have thick trunks and dense green foliage, creating a shaded canopy over the path. The ground is covered with fallen leaves and some low-lying plants in wooden planters.

**Why are branches  
"cheap"?**

# Branching (structure) comes for free

- Together, commits and parent relations form the git history DAG<sup>6</sup>
- Multiple commits can share a parent => natural "branching" structure
- Could theoretically manage divergent version paths without an explicit "branch" concept<sup>7</sup>

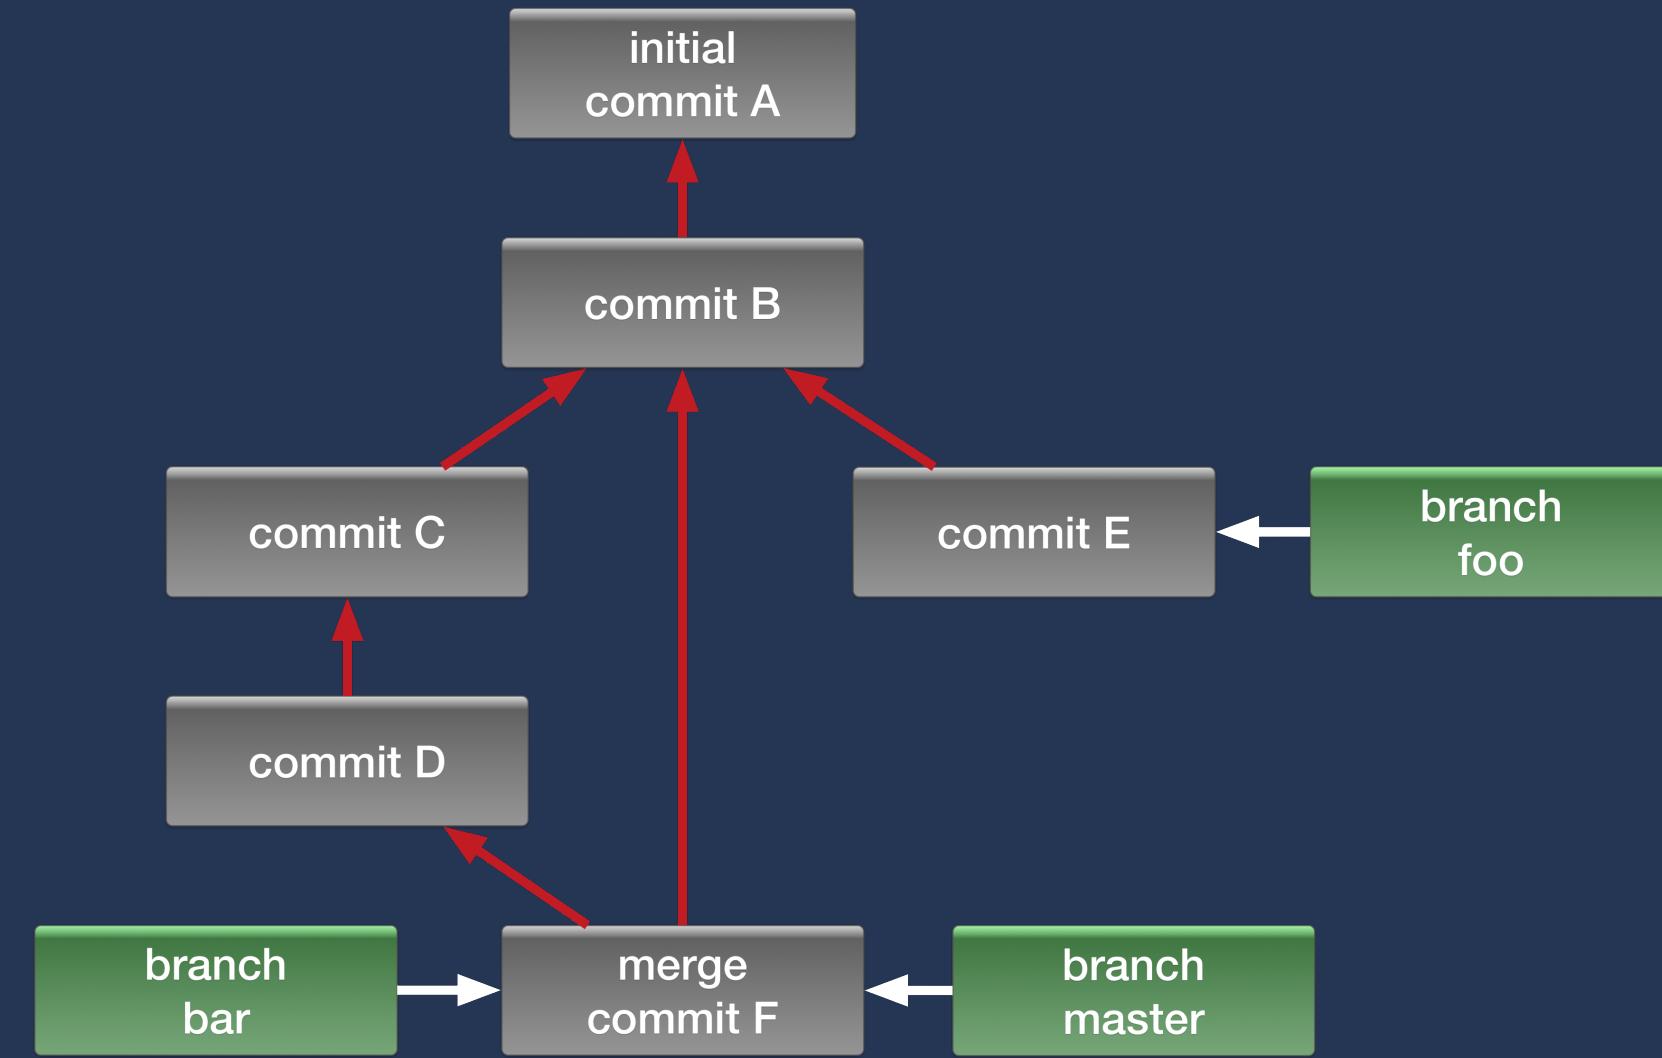


<sup>6</sup> It can be further specified as a rooted connected directed acyclic graph. 😮  
Note that the history is *not* a tree because commits can have multiple parents, but it is tree-like in other respects.

<sup>7</sup> It would involve manually tracking commit SHAs, though. 🤦

# A git branch (object) is just a pointer

- Git's "branch" object (stored as reference to a commit SHA) affords two major conveniences:
  - Nice name for checkouts, etc
  - The checked-out branch moves forward with each new commit<sup>8</sup>
- Note: there is nothing special about **master**: it's a regular branch<sup>9</sup>



<sup>8</sup> Unlike tags (similarly just pointers), which stay put unless explicitly moved.

<sup>9</sup> The branch created by `git init` is called "master" by default.

# Ergo, branches are cheap

- Creating a branch == creating a SHA reference: cheap!
- Because git only creates new file snapshots for modified files, they are also cheap to maintain<sup>10</sup>
- Deleting a branch only deletes the reference: also cheap!
- Bonus: the commits still exist and can be recovered

---

<sup>10</sup> Relative to other VCSs that maintain an entirely separate project copy per branch.

# Merges are (fairly) easy

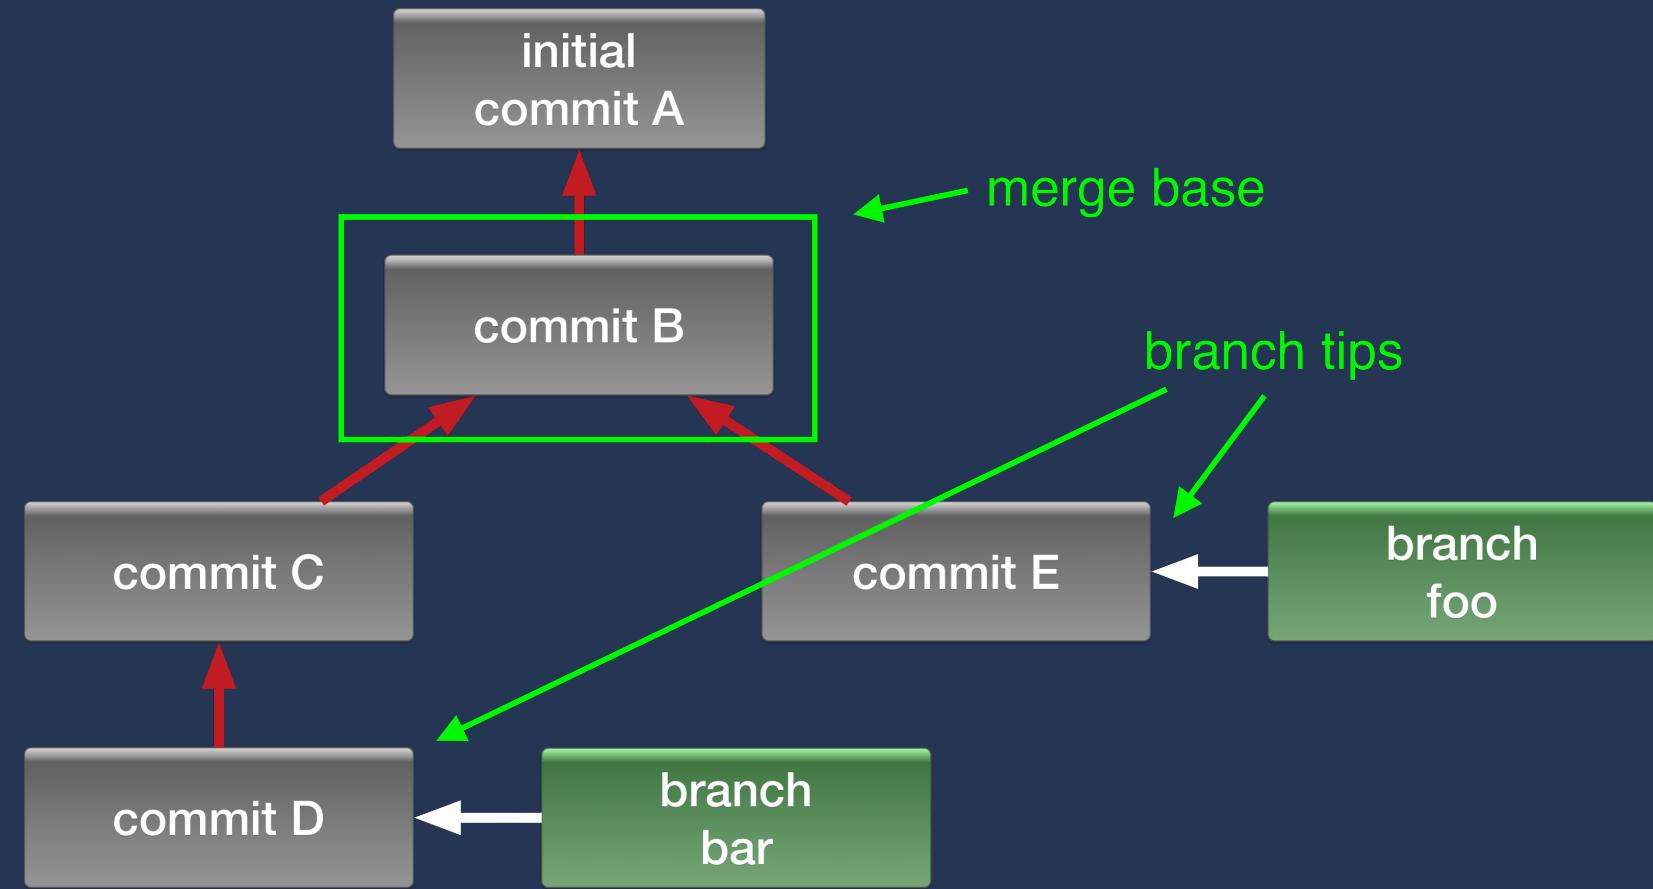
- To merge, git compares branches to their best **merge base**
- The merge base (most recent common ancestor) is easily determined from the commit graph
- Unlike a simple 3-point merge, git preserves granular history info by replaying commits from one branch onto the other
- This allows git to correctly handle many tricky merge

# Example merge scenario

To merge bar into foo:

```
$ git checkout foo  
$ git merge bar
```

- Determine merge base
- Compute diffs ( $C - B$ ) and ( $D - C$ )
- Apply diffs in order onto E
- Turn the result into a merge commit
- Move branch foo to merge commit

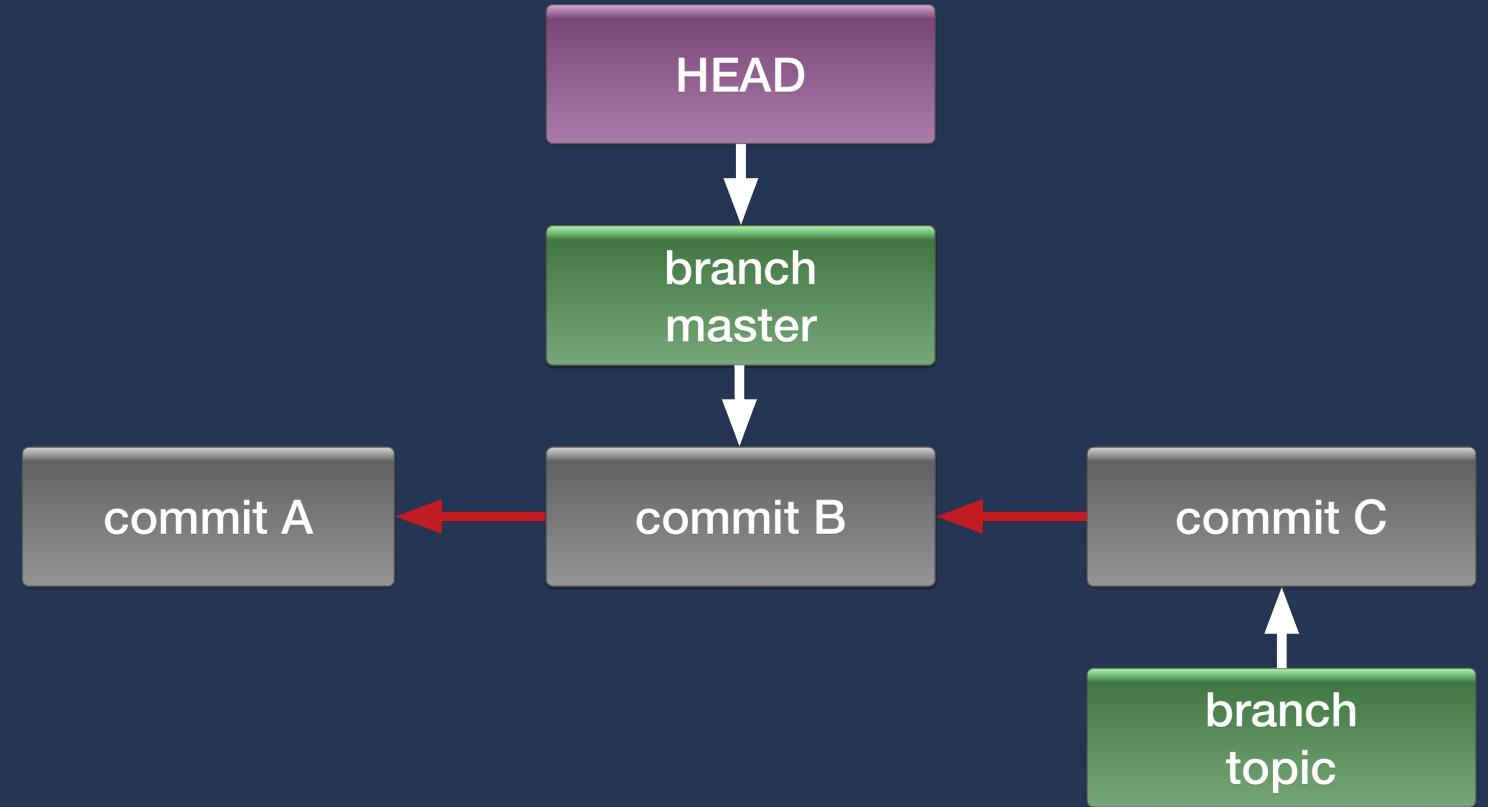


A scenic view of a park with large trees, a path, and a building in the background.

# How do checkouts work?

# Checkouts: HEAD

- The HEAD reference determines "where you are" in the commit graph
- HEAD can point either to a branch reference or directly to a commit<sup>11</sup>
- Example (right): The master branch is currently "checked out"



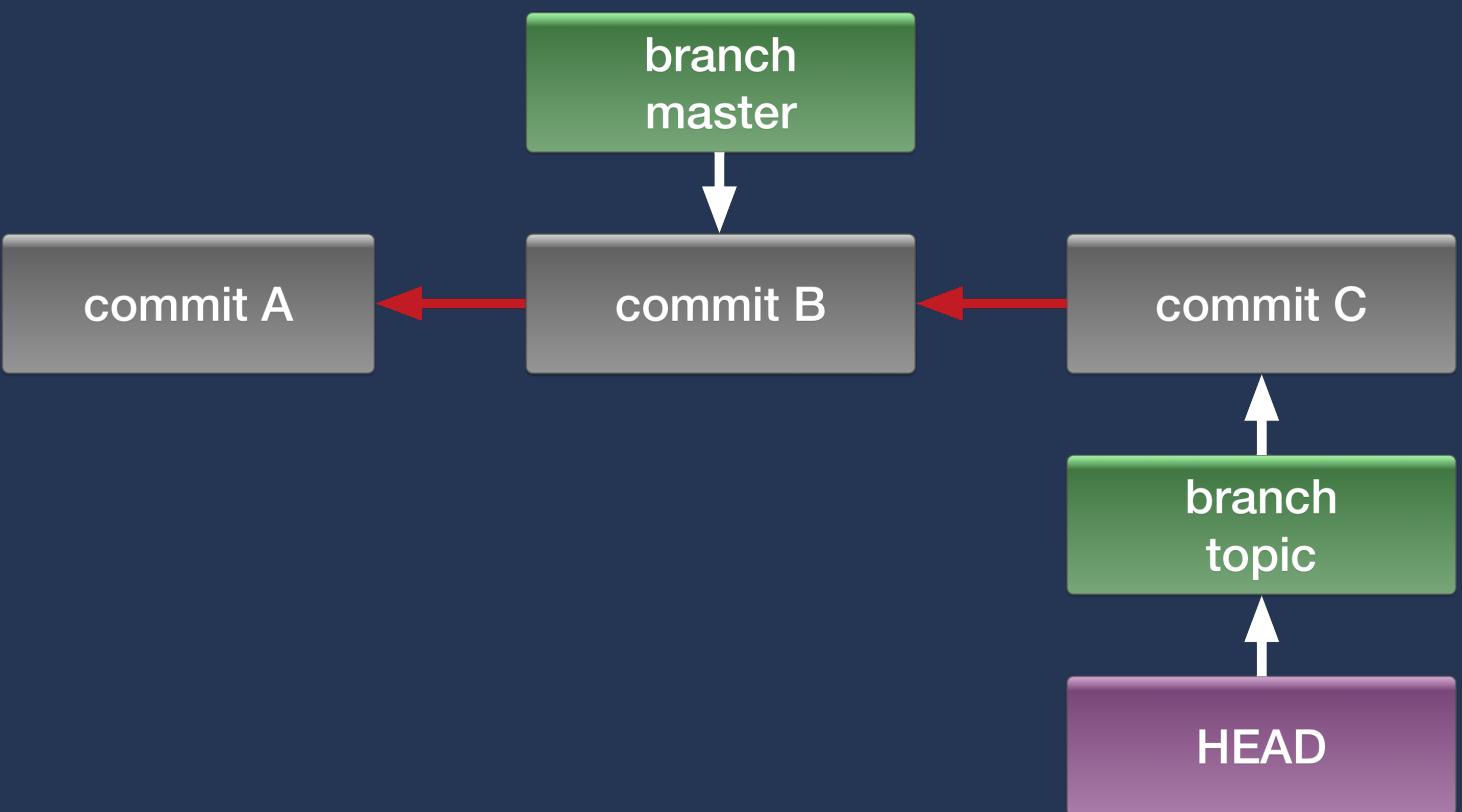
---

<sup>11</sup> This is the "unattached HEAD" state.

# Checkouts: Switching branches

```
$ git checkout topic
```

- Modify HEAD to point to topic
- Copy commit C's snapshot tree to the staging area
- Decompress the files in the project snapshot and copy them to the working directory<sup>12</sup>

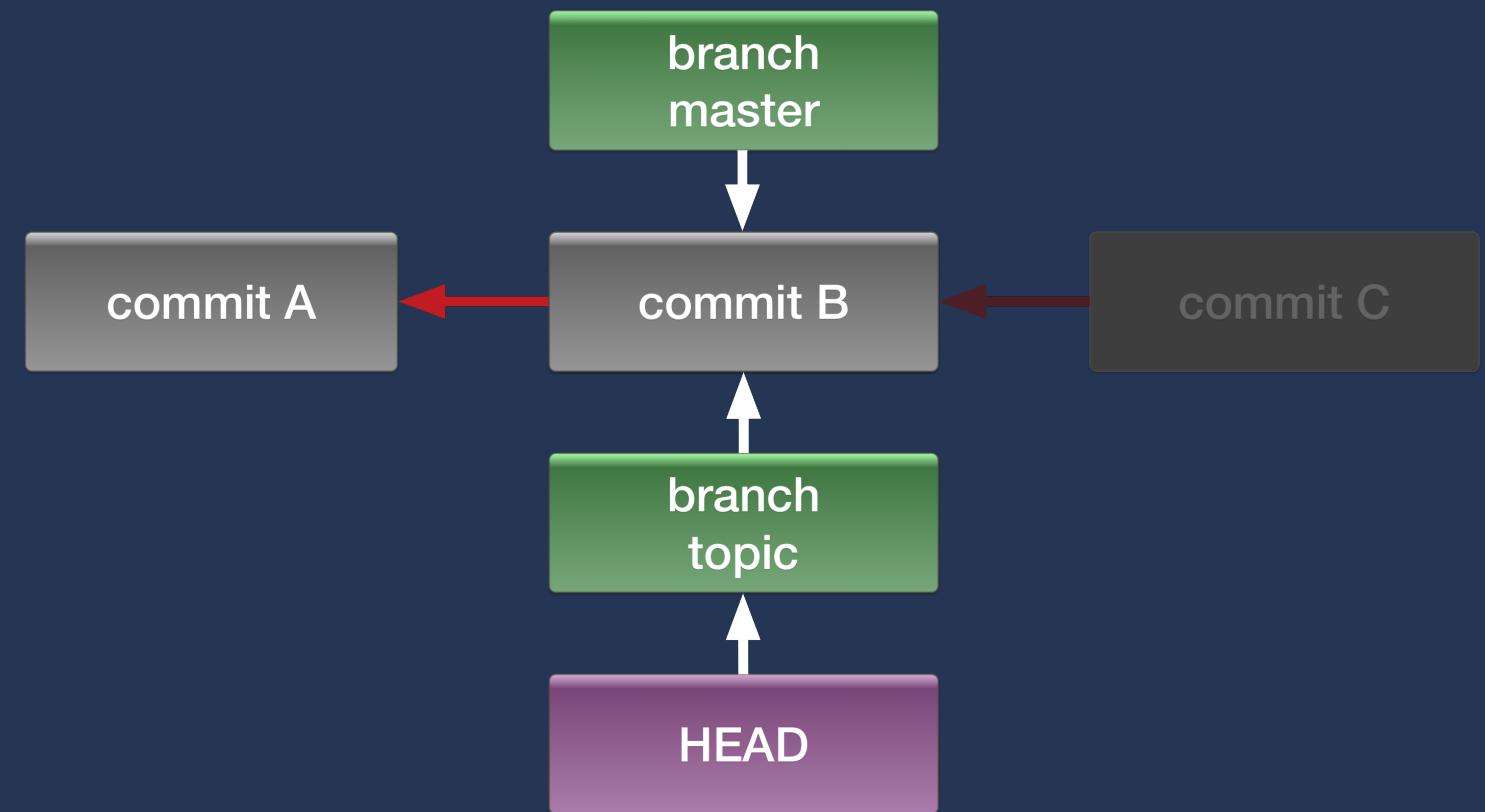


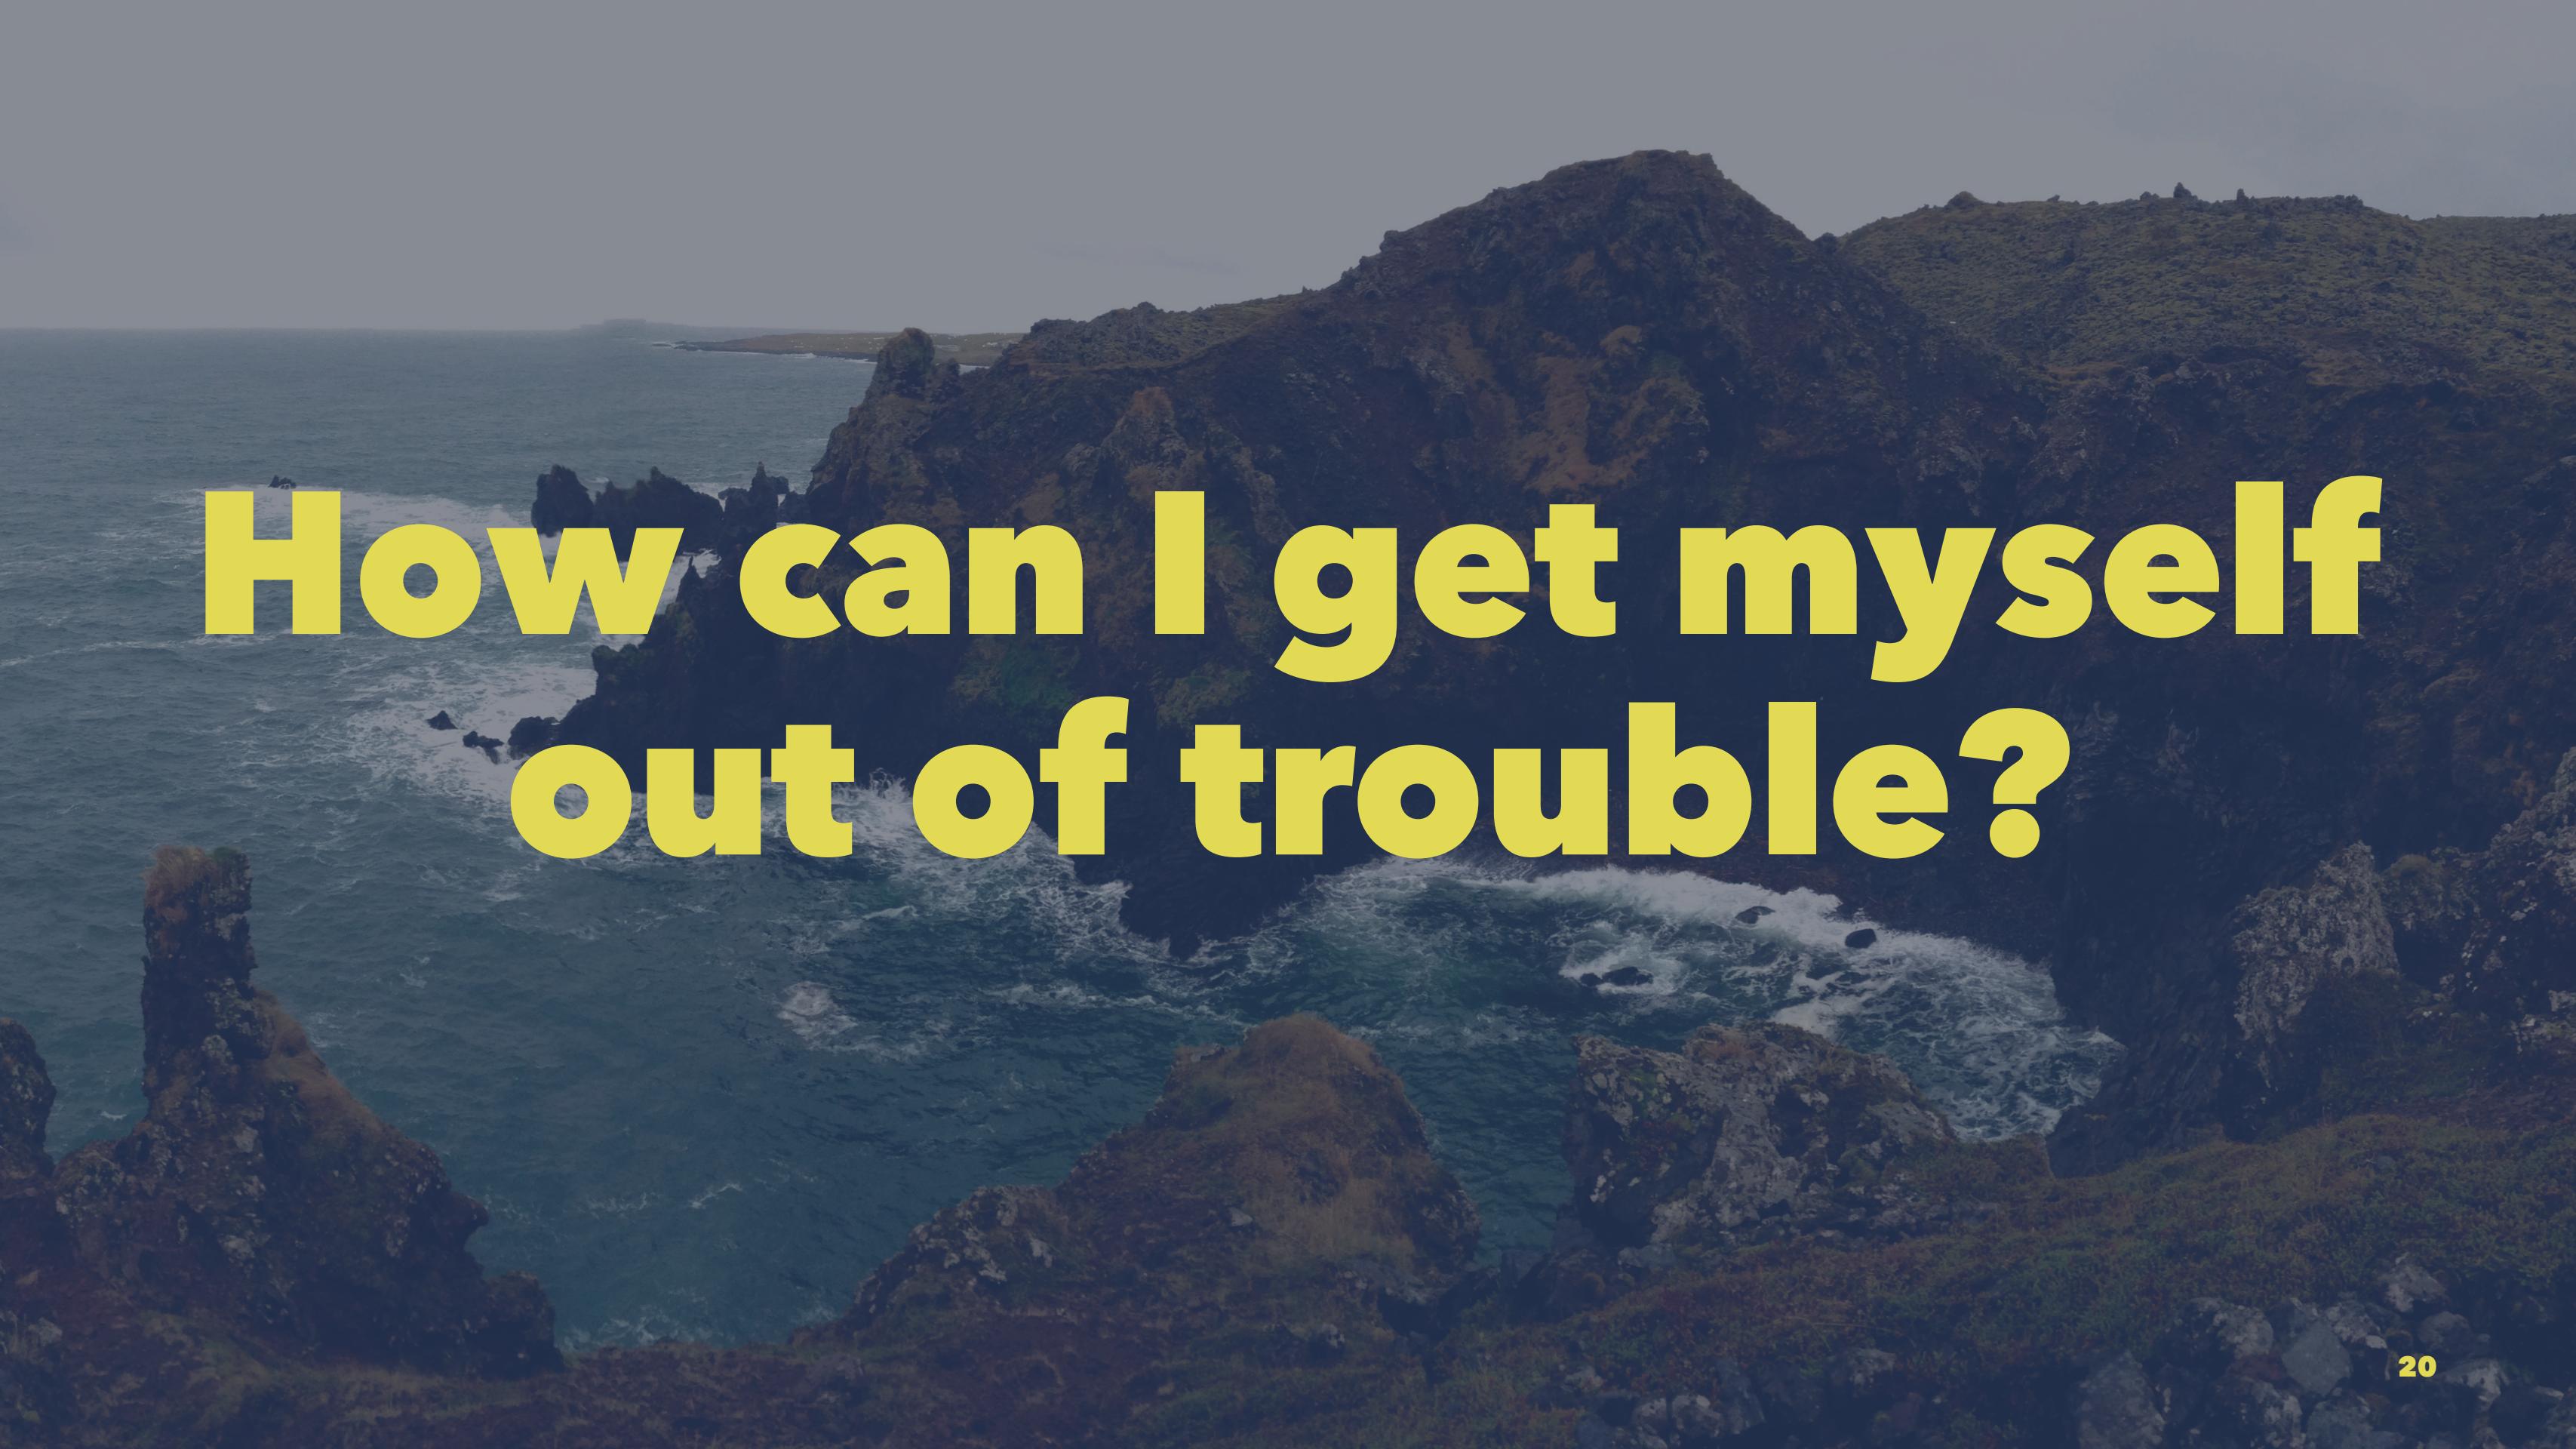
<sup>12</sup> This could clobber uncommitted changes in your working directory, which is why git may throw an error if you try to do a checkout with a dirty working directory.

# Resets are like checkouts

```
git reset --<mode> master
```

- A **hard** reset does the same 3 steps as a checkout, except that the pointer that moves is the **branch**, rather than HEAD
- A default (mode=**mixed**) reset skips the working directory overwrite
- A **soft** reset also skips the staging area overwrite



A dark, moody landscape of a rugged coastline. In the foreground, dark, jagged rock formations rise from a turbulent, blue-green ocean. The water is filled with white-capped waves crashing against the rocks. In the background, a large, dark mountain or cliff face rises, its slopes covered in sparse vegetation. The sky is a uniform, dark grey, suggesting overcast weather.

**How can I get myself  
out of trouble?**

# Meet the reflog

## Your new best friend

- The reflog is a **local-only** log of all changes to git **refs**, including branches, tags, HEAD, stashes
- By default, git reflog shows you a log of HEAD changes
- git reflog <ref name> to view changes to another ref
- The reflog can be used to **return to a previous state**<sup>13</sup>

---

<sup>13</sup> A previous **committed** state. If you accidentally deleted uncommitted work, no dice. Commit early and often!

# What can I find in the reflog?

- Some of the changes recorded in the reflog:
  - new commits (including merge commits, cherry-picks)
  - modifications to commits
  - branch or commit checkouts
- Fetches or pushes to a remote are not recorded in the reflog, because they don't affect your local repository copy

# Usecase

## Roll back to a previous state

- Use the reflog to immediately roll back from a git mistake (e.g. botched rebase, pulled instead of fetched)<sup>14</sup>
- Identify the HEAD reference **before** the error, then do a hard reset to it
- Ex: `git reset --hard HEAD@{1}`

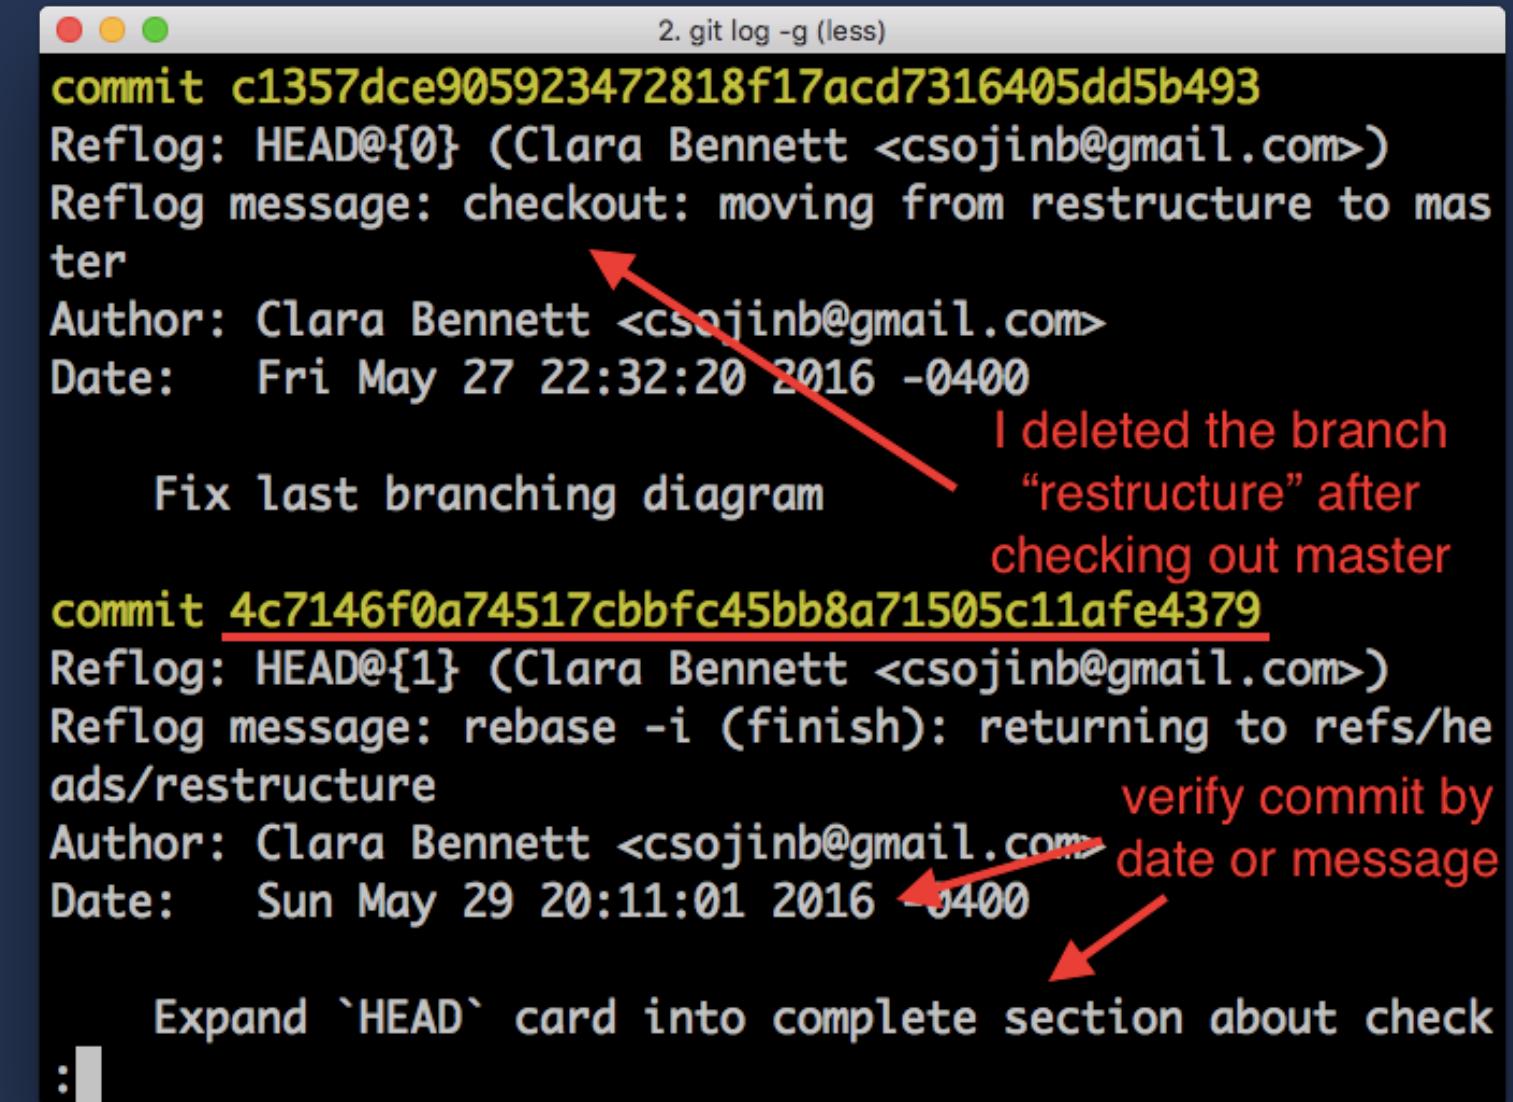
```
2. git reflog (git)
19211d3 HEAD@{0}: reset: moving to 19211d3 undo please!
217dac1 HEAD@{1}: rebase -i (finish): returning to refs/
heads/restructure
217dac1 HEAD@{2}: rebase -i (pick): Expand `HEAD` card i
nto complete section about checkouts and resets
4cb6e58 HEAD@{3}: rebase -i (pick): Restructure branchin
g section
19211d3 HEAD@{4}: rebase -i (pick): Improve diagram reso
lution
7f501c3 HEAD@{5}: rebase -i (pick): Restructure commit s
torage section
:
```

<sup>14</sup> You can even use this to recover from a bad reflog reset!

# Usecase

## Recover a deleted branch

- We can't use the branch-specific reflog because it was deleted too
- View detailed commit information in the HEAD log with `git log -g`
- Find the SHA of the former branch tip and remake the branch:  
`$ git branch recovery 4c7146f`
- This technique can also be used to recover modified commits



```
2. git log -g (less)
commit c1357dce905923472818f17acd7316405dd5b493
Reflog: HEAD@{0} (Clara Bennett <csojinb@gmail.com>)
Reflog message: checkout: moving from restructure to master
Author: Clara Bennett <csojinb@gmail.com>
Date:   Fri May 27 22:32:20 2016 -0400
      Fix last branching diagram
commit 4c7146f0a74517cbbfc45bb8a71505c11afe4379
Reflog: HEAD@{1} (Clara Bennett <csojinb@gmail.com>)
Reflog message: rebase -i (finish): returning to refs/heads/restructure
Author: Clara Bennett <csojinb@gmail.com>
Date:   Sun May 29 20:11:01 2016 -0400
      verify commit by date or message
      Expand `HEAD` card into complete section about check
:|
```

# Off-branch commits aren't stored forever

- Git is conservative: it keeps commits reachable by any reference, including the reflog
- Default reflog expire time is 90 days
- Unless you explicitly trigger garbage collection, "expired" reflog items are only cleaned up if there's a space issue<sup>15</sup>
- With the defaults, reflog expiry unlikely to cause issues

---

<sup>15</sup> So, a small repo that only you contribute to could still have commits from old branches from a year ago, for example.

# What next?

- Go forth and git greatly!
- This presentation can be found at  
[github.com/csojinb/git-under-the-hood](https://github.com/csojinb/git-under-the-hood)
- Scott Chacon's book Pro Git (free!) is an excellent resource
  - To learn more about git internals in particular, check out Chapter 10 and take a swim through your .git directory



@csojinb